

Snow64 Instruction Set

- Notes
 - There are no Instruction LARs (the typical instruction fetch of most computer processors is used instead).
 - Addresses are 64-bit.
- Registers
 - The DLARs themselves (there are 16, but this may be changed later):
 - dzero (always zero),
 - du0, du1, du2, du3 du4, du5, du6, du7, du8, du9, du10, du11 (user registers)
 - dlr (standard link register (hardware does not enforce this))
 - dfp (standard frame pointer (hardware does not enforce this))
 - dsp (standard stack pointer (hardware does not enforce this))
 - Other registers:
 - pc (the program counter, 64-bit)

Instruction Set

- Note: All invalid instructions are treated as NOPs.
- ALU Instructions: Opcode Group: 0b000
 - Encoding: 000t aaaa bbbb cccc oooo iiii iiii iiii
 - t: operation type: if 0b0: scalar operation; else: vector operation
 - a: dDest
 - b: dSrc0
 - c: dSrc1
 - o: opcode
 - i: 12-bit signed immediate
 - Note: For ALU instructions, any result that doesn't fit in the destination will be truncated to fit into the destination. This affects both scalar and vector operations.
 - Note: also, for each of these instructions, the address field is not used as an operand, just the data field.
 - Example: adds d1, d2, d3
 - Effect: $d1.sdata \leftarrow \text{cast_to_type_of_d1}(d2.sdata) + \text{cast_to_type_of_d1}(d3.sdata)$
 - Instructions:
 - **add** dDest, dSrc0, dSrc1
 - Opcode: 0x0
 - Scalar Mnemonic: adds
 - Vector Mnemonic: addv
 - **sub** dDest, dSrc0, dSrc1
 - Opcode: 0x1
 - Scalar Mnemonic: subs
 - Vector Mnemonic: subv

- **slt** dDest, dSrc0, dSrc1
 - Opcode: 0x2
 - Scalar Mnemonic: slts
 - Vector Mnemonic: sltv
 - Note: set less than
 - Note: The signedness of dDest will be used for the operation
- **mul** dDest, dSrc0, dSrc1
 - Opcode: 0x3
 - Scalar Mnemonic: muls
 - Vector Mnemonic: mulv
 - Note: If dDest has a larger size than both dSrc0 and dSrc1, then the signedness used for the operation will be that of dDest
 - Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
- **div** dDest, dSrc0, dSrc1
 - Opcode: 0x4
 - Scalar Mnemonic: divs
 - Vector Mnemonic: divv
 - Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
- **and** dDest, dSrc0, dSrc1
 - Opcode: 0x5
 - Scalar Mnemonic: ands
 - Vector Mnemonic: andv
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **orr** dDest, dSrc0, dSrc1
 - Opcode: 0x6
 - Scalar Mnemonic: orrs
 - Vector Mnemonic: orrv
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **xor** dDest, dSrc0, dSrc1
 - Opcode: 0x7
 - Scalar Mnemonic: xors
 - Vector Mnemonic: xorv
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **shl** dDest, dSrc0, dSrc1
 - Opcode: 0x8
 - Scalar Mnemonic: shls
 - Vector Mnemonic: shl v
 - Note: Shift left
 - Note: this operation always uses 64-bit components of the input register(s) (no casting is performed), and saves the result to the destination register as if the destination register was tagged as 64-bit
- **shr** dDest, dSrc0, dSrc1
 - Opcode: 0x9
 - Scalar Mnemonic: shrs

- Vector Mnemonic: shrv
 - Note: Shift right
 - Note: dSrc0's signedness is used to determine the type of right shift:
 - If dSrc0 is unsigned, a logical right shift is performed
 - If dSrc0 is signed, an arithmetic right shift is performed
 - Note: dSrc1 is always treated as unsigned (due to being a number of bits to shift by)
 - Note: this operation always uses 64-bit components of the input register(s) (no casting is performed), and saves the result to the destination register as if the destination register was tagged as 64-bit
 - **inv** dDest, dSrc0
 - Opcode: 0xa
 - Scalar Mnemonic: invs
 - Vector Mnemonic: invv
 - Note: Bitwise invert
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **not** dDest, dSrc0
 - Opcode: 0xb
 - Scalar Mnemonic: nots
 - Vector Mnemonic: notv
 - Note: Logical not
 - Note: this operation always uses 64-bit components of the input register(s) (no casting is performed), and saves the result to the destination register as if the destination register was tagged as 64-bit
 - **addi** dDest, pc, simm12
 - Opcode: 0xc
 - Scalar Mnemonic: addis
 - Vector Mnemonic: addiv
 - Note: This is useful for pc-relative loads, relative branches, and for getting the return address of a subroutine call into a LAR before jumping to a subroutine.
 - **addi** dDest, dSrc0, simm12
 - Opcode: 0xd
 - Scalar Mnemonic: addis
 - Vector Mnemonic: addiv
 - Note: This instruction can operate as "load immediate" by using dzero as dSrc0, but note that this instruction does not affect the dDest.address field.
- Instructions for interacting with special-purpose registers:
 Opcode Group: 0b001
 - Encoding: 0010 aaaa oooo iiii iiii iiii iiii iiii
 - a: dA
 - o: opcode
 - i: 20-bit signed immediate

- Note: all instructions in group 0b001 are scalar operations.
- Note: dX.sdata is simply the current scalar portion of the data LAR called dX
- Instructions:
 - **btru** dA, simm20
 - Opcode: 0x0
 - Effect: if (dA.sdata != 0) pc <= pc + sign_extend_to_64(simm20);
 - **bfal** dA, simm20
 - Opcode: 0x1
 - Effect: if (dA.sdata == 0) pc <= pc + sign_extend_to_64(simm20);
 - **jmp** dA
 - Opcode: 0x2
 - Effect: pc <= dA.sdata;
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
- Load Instructions: Opcode Group: 0b010
 - Encoding: 0100 aaaa bbbb cccc oooo iiii iiii iiii
 - a: dDest
 - b: dSrc0
 - c: dSrc1
 - o: opcode
 - i: 12-bit signed immediate
 - Effect:
 - Load LAR-sized data from 64-bit address computed as follows: (dB.address + extend_to_64(dC.sdata) + (sign_extend_to_64(simm12)))
 - This 64-bit address is referred to as the "effective address".
 - The type of extension of the extend_to_64(dC.sdata) expression is based upon the type of dC.
 - If dC is tagged as an unsigned integer, zero-extension is performed.
 - If dC is tagged as a signed integer, sign-extension is performed.
 - If dC is tagged as a BFloat16, dC.sdata is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
 - Due to associativity of the LARs, these instructions will not actually load from memory if the effective address's data already loaded into a LAR.
 - Instructions:
 - **ldu8** dA, dB, dC, simm12
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **lds8** dA, dB, dC, simm12
 - Opcode: 0x1
 - Note: signed 8-bit integer(s)

- **ldu16** dA, dB, dC, simm12
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
 - **lds16** dA, dB, dC, simm12
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
 - **ldu32** dA, dB, dC, simm12
 - Opcode: 0x4
 - Note: unsigned 32-bit integer(s)
 - **lds32** dA, dB, dC, simm12
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
 - **ldu64** dA, dB, dC, simm12
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
 - **lds64** dA, dB, dC, simm12
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
 - **ldf16** dA, dB, dC, simm12
 - Opcode: 0x8
 - Note: BFloat16 format floating point number.
- Store Instructions: Opcode Group: 0b011
 - Encoding: 0110 aaaa bbbb cccc oooo iiii iiii iiii
 - a: dA
 - b: dB
 - c: dC
 - o: opcode
 - i: 12-bit signed immediate
 - Note: These are actually type conversion instructions as actual writes to memory are done lazily
 - Effect:
 - These instructions mark dA as dirty, change its address to the effective address (see next bullet), and sets its type.
 - The 64-bit effective address is computed as follows:
 $(dB.address + extend_to_64(dC.sdata) + (sign_extend_to_64(simm12)))$
 - The type of extension of the $extend_to_64(dC.sdata)$ expression is based upon the type of dC.
 - If dC is tagged as an unsigned integer, zero-extension is performed.
 - If dC is tagged as a signed integer, sign-extension is performed.
 - If dC is tagged as a BFloat16, dC.sdata is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
 - Instructions:
 - **stu8** dA, dB, dC, simm12
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **sts8** dA, dB, dC, simm12
 - Opcode: 0x1

- Note: signed 8-bit integer(s)
- **stu16** dA, dB, dC, simm12
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
- **sts16** dA, dB, dC, simm12
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
- **stu32** dA, dB, dC, simm12
 - Opcode: 0x4
 - Note: unsigned 32-bit integer(s)
- **sts32** dA, dB, dC, simm12
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
- **stu64** dA, dB, dC, simm12
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
- **sts64** dA, dB, dC, simm12
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
- **stf16** dA, dB, dC, simm12
 - Opcode: 0x8
 - Note: BFloat16 format floating point number.