

Snow64 Instruction Set

- Notes
 - There are no Instruction LARs (the typical instruction fetch of most computer processors is used instead).
 - Addresses are 64-bit.
 - Interrupts are supported (this may be a first for a LARs architecture)
 - Port-mapped I/O is supported (this may be a first for a LARs architecture)
- Data LARs:

```
typedef struct packed
{
    // Data field
    union packed
    {
        // This is possibly not valid SystemVerilog because of
        // arrays inside a packed struct, but it makes for nice
        // pseudocode
        logic [7:0] data_8[0:31];

        logic [15:0] data_16[0:15];

        struct packed
        {
            // sign bit, 1 means negative
            logic sign;

            // Exponent, +bias
            logic [7:0] exp;

            // Mantissa; normalized implies 1 MSB
            logic [6:0] mant;
        } data_float_16[0:15];

        logic [31:0] data_32[0:7];
        logic [63:0] data_64[0:3];
    } data;

    // Note that this is a 64-bit structure
    union packed
    {
        struct packed
        {
            logic [63 - 5 : 0] base_ptr;

            logic [4:0] offset;
        } addr_8;

        // Used for both 16-bit integers and the half floats
        struct packed
        {
            logic [63 - 4 : 0] base_ptr;

            logic [3:0] offset;
        } addr_16;

        struct packed
        {
```

```

        logic [63 - 3 : 0] base_ptr;

        logic [2:0] offset;
    } addr_32;

    struct packed
    {
        logic [63 - 2 : 0] base_ptr;

        logic [1:0] offset;
    } addr_64;
} addr;

// Integer Type (used when "is_float_16" is 1'b0):
// 2'b00: 8-bit
// 2'b01: 16-bit
// 2'b10: 32-bit
// 2'b11: 64-bit
logic [1:0] type_of_int;

// These are actually packed 16-bit floats, the implementing the
// high 16 bits of 32-bit IEEE float.
// "type_of_int" is ignored when "is_float_16" is 1'b1.
logic is_float_16;

// Unsigned: 1'b0
// Signed: 1'b1
// Note: unsgn_or_sgn is ignored for floats
logic unsgn_or_sgn;

// Data should be lazily stored to memory if this is 1'b1
// Otherwise, when this is 1'b0, data in this LAR is up to date
// with memory.
logic dirty;

} DataLar;

```

- Registers

- The DLARs themselves (there are 16, but this may be changed later):
 - `dzero` (always zero),
 - `du0` , `du1` , `du2` , `du3` , `du4` , `du5` , `du6` , `du7` , `du8` , `du9` , `du10` , `du11` (user registers)
 - `d1r` (standard link register (hardware does not enforce this))
 - `dfp` (standard frame pointer (hardware does not enforce this))
 - `dsp` (standard stack pointer (hardware does not enforce this))
- Other registers:
 - `pc` (the program counter, 64-bit)
 - `ie` (whether or not interrupts are enabled, 1-bit)
 - `ireta` (the interrupt return address, 64-bit)
 - `idsta` (the interrupt destination address, 64-bit; upon an interrupt, the program counter is set to the value in this register)

Instruction Set

- Note: All invalid instructions are treated as NOPs.
- ALU Instructions: Opcode Group: 0b000
 - Encoding: `000t aaaa bbbb cccc oooo iiii iiii iiii`
 - `t` : operation type: if `0b0` : scalar operation; else: vector operation
 - `a` : dDest
 - `b` : dSrc0

- `c` : dSrc1
- `o` : opcode
- `i` : 12-bit signed immediate
- Note: For ALU instructions, any result that doesn't fit in the destination will be truncated to fit into the destination. This affects both scalar and vector operations.
- Note: also, for each of these instructions, the address field is not used as an operand, just the data field.
- Instructions:
 - **add** dDest, dSrc0, dSrc1
 - Opcode: 0x0
 - Scalar Mnemonic: `adds`
 - Vector Mnemonic: `addv`
 - **sub** dDest, dSrc0, dSrc1
 - Opcode: 0x1
 - Scalar Mnemonic: `subs`
 - Vector Mnemonic: `subv`
 - **slt** dDest, dSrc0, dSrc1
 - Opcode: 0x2
 - Scalar Mnemonic: `slts`
 - Vector Mnemonic: `sltv`
 - Note: set less than
 - Note: The signedness of dDest will be used for the operation
 - **mul** dDest, dSrc0, dSrc1
 - Opcode: 0x3
 - Scalar Mnemonic: `muls`
 - Vector Mnemonic: `mulv`
 - Note: If dDest has a larger size than both dSrc0 and dSrc1, then the signedness used for the operation will be that of dDest
 - Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
 - **div** dDest, dSrc0, dSrc1
 - Opcode: 0x4
 - Scalar Mnemonic: `divs`
 - Vector Mnemonic: `divv`
 - Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
 - **and** dDest, dSrc0, dSrc1
 - Opcode: 0x5
 - Scalar Mnemonic: `ands`
 - Vector Mnemonic: `andv`
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **orr** dDest, dSrc0, dSrc1
 - Opcode: 0x6
 - Scalar Mnemonic: `orrs`
 - Vector Mnemonic: `orrv`
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **xor** dDest, dSrc0, dSrc1
 - Opcode: 0x7
 - Scalar Mnemonic: `xors`
 - Vector Mnemonic: `xorv`
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **shl** dDest, dSrc0, dSrc1
 - Opcode: 0x8
 - Scalar Mnemonic: `shls`
 - Vector Mnemonic: `shlv`
 - Note: Shift left
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **shr** dDest, dSrc0, dSrc1

- Opcode: 0x9
 - Scalar Mnemonic: `shrs`
 - Vector Mnemonic: `shrv`
 - Note: Shift right
 - Note: dSrc0's signedness is used to determine the type of right shift:
 - If dSrc0 is unsigned, a logic right shift is performed
 - If dSrc0 is signed, an arithmetic right shift is performed
 - Note: dSrc1 is always treated as unsigned (due to being a number of bits to shift by)
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **inv** dDest, dSrc0
 - Opcode: 0xa
 - Scalar Mnemonic: `invs`
 - Vector Mnemonic: `invv`
 - Note: Bitwise invert
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **not** dDest, dSrc0
 - Opcode: 0xb
 - Scalar Mnemonic: `nots`
 - Vector Mnemonic: `notv`
 - Note: Logical not
 - **add** dDest, pc, simm12
 - Opcode: 0xc
 - Scalar Mnemonic: `adds`
 - Vector Mnemonic: `addv`
 - Note: This is useful for pc-relative loads, relative branches, and for getting the return address of a subroutine call into a LAR before jumping to a subroutine.
- Instructions for interacting with special-purpose registers:
 - Opcode Group: 0b001
 - Encoding: `0010 aaaa oooo iiii iiii iiii iiii iiii`
 - `a`: dA
 - `o`: opcode
 - `i`: 20-bit signed immediate
 - Note: all instructions in group 0b001 are scalar operations.
 - Note: `dx.sdata` is simply the current scalar portion of the data LAR called `dx`
 - Instructions:
 - **btru** dA, simm20
 - Opcode: 0x0
 - Effect: `if (dA.sdata != 0) pc <= pc + sign_extend_to_64(simm20);`
 - **bfal** dA, simm20
 - Opcode: 0x1
 - Effect: `if (dA.sdata == 0) pc <= pc + sign_extend_to_64(simm20);`
 - **jmp** dA
 - Opcode: 0x2
 - Effect: `pc <= dA.sdata;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **ei**
 - Opcode: 0x3
 - Effect: `ie <= 1'b1;`
 - Note: Enable interrupts
 - **di**
 - Opcode: 0x4
 - Effect: `ie <= 1'b0;`
 - Note: Disable interrupts
 - **reti**

- Opcode: 0x5
 - Effect: `ie <= 1'b1; pc <= ireta;`
 - Note: Return from an interrupt
 - **cpy** dA, ie
 - Opcode: 0x6
 - Effect: `dA.sdata <= ie; // acts differently if dA is tagged as a float`
 - **cpy** dA, ireta
 - Opcode: 0x7
 - Effect: `dA.sdata <= ireta;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **cpy** dA, idsta
 - Opcode: 0x8
 - Effect: `dA.sdata <= idsta;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **cpy** ie, dA
 - Opcode: 0x9
 - Effect: `ie <= (dA.sdata != 0);`
 - **cpy** ireta, dA
 - Opcode: 0xa
 - Effect: `ireta <= dA.sdata;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **cpy** idsta, dA
 - Opcode: 0xb
 - Effect: `idsta <= dA.sdata;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
- Load Instructions: Opcode Group: 0b010
 - Encoding: `0100 aaaa bbbb cccc oooo iiii iiii iiii`
 - `a` : dDest
 - `b` : dSrc0
 - `c` : dSrc1
 - `o` : opcode
 - `i` : 12-bit signed immediate
 - Effect:
 - Load LAR-sized data from 64-bit address computed as follows: `(dB.address + extend_to_64(dC.sdata) + (sign_extend_to_64(simm12)))`
 - This 64-bit address is referred to as the "effective address".
 - The type of extension of the `extend_to_64(dC.sdata)` expression is based upon the type of `dC` .
 - If `dC` is tagged as an unsigned integer, zero-extension is performed.
 - If `dC` is tagged as a signed integer, sign-extension is performed.
 - If `dC` is tagged as a BFloat16, `dC.sdata` is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
 - Due to associativity of the LARs, these instructions will not actually load from memory if the effective address's data already loaded into a LAR.
 - Instructions:
 - **ldu8** dA, dB, dC, simm12
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **lds8** dA, dB, dC, simm12
 - Opcode: 0x1
 - Note: signed 8-bit integer(s)

- **ldu16** dA, dB, dC, simm12
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
 - **lds16** dA, dB, dC, simm12
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
 - **ldu32** dA, dB, dC, simm12
 - Opcode: 0x4
 - Note: unsigned 32-bit integer(s)
 - **lds32** dA, dB, dC, simm12
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
 - **ldu64** dA, dB, dC, simm12
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
 - **lds64** dA, dB, dC, simm12
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
 - **ldf16** dA, dB, dC, simm12
 - Opcode: 0x8
 - Note: BFloat16 format floating point number.
- Store Instructions: Opcode Group: 0b011
 - Encoding: 0110 aaaa bbbb cccc oooo iiii iiii iiii
 - a : dA
 - b : dB
 - c : dC
 - o : opcode
 - i : 12-bit signed immediate
 - Note: These are actually type conversion instructions as actual writes to memory are done lazily
 - Effect:
 - These instructions marks dA as dirty, change its address to the effective address (see next bullet), and sets its type.
 - The 64-bit effective address is computed as follows: $(dB.address + extend_to_64(dC.sdata) + (sign_extend_to_64(simm12)))$
 - The type of extension of the `extend_to_64(dC.sdata)` expression is based upon the type of dC .
 - If dC is tagged as an unsigned integer, zero-extension is performed.
 - If dC is tagged as a signed integer, sign-extension is performed.
 - If dC is tagged as a BFloat16, dC.sdata is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
 - Instructions:
 - **stu8** dA, dB, dC, simm12
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **sts8** dA, dB, dC, simm12
 - Opcode: 0x1
 - Note: signed 8-bit integer(s)
 - **stu16** dA, dB, dC, simm12
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
 - **sts16** dA, dB, dC, simm12
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
 - **stu32** dA, dB, dC, simm12
 - Opcode: 0x4

- Note: unsigned 32-bit integer(s)
 - **sts32** dA, dB, dC, simm12
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
 - **stu64** dA, dB, dC, simm12
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
 - **sts64** dA, dB, dC, simm12
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
 - **stf16** dA, dB, dC, simm12
 - Opcode: 0x8
 - Note: BFloat16 format floating point number.
- Port-mapped Input/Output Instructions: Opcode Group: 0b100
 - Encoding: 100t aaaa bbbb oooo iiii iiii iiii iiii
 - t : operation type: if 0b0 : scalar operation; else: vector operation
 - a : dA
 - b : dB
 - o : opcode
 - i : 16-bit signed immediate
 - Note: dx.sdata is simply the current scalar portion of the data LAR called dx
 - Note: For the in... instructions, the entirety of dA.data is set to the received data. The type of dA is set based upon the instruction opcode.
 - Note: For outs, dA.sdata is sent to the output port, along with the type of data (in case the particular I/O port cares).
 - Note: For outv, the entirety of dA.data is sent to the output port, along with the type of data (in case the particular I/O port cares).
 - Note: For each of these instructions, the I/O address used is computed by the formula

$$\text{cast_to_64}(\text{dB.sdata}) + \text{sign_extend_to_64}(\text{simm16})$$
 - Instructions:
 - **inu8** dA, dB, simm16
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **ins8** dA, dB, simm16
 - Opcode: 0x1
 - Note: signed 8-bit integer(s)
 - **inu16** dA, dB, simm16
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
 - **ins16** dA, dB, simm16
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
 - **inu32** dA, dB, simm16
 - Opcode: 0x4
 - Note: unsigned 32-bit integer(s)
 - **ins32** dA, dB, simm16
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
 - **inu64** dA, dB, simm16
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
 - **ins64** dA, dB, simm16
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
 - **inf16** dA, dB, simm16

- Opcode: 0x8
- Note: BFloat16 format floating point number.
- **out** (actual mnemonics below)
 - Opcode: 0x9
 - **outs** dA, dB, simm16
 - t : 0
 - Note: dA.sdata is simply sent to the output data bus.
 - **outv** dA, dB, simm16
 - t : 1
 - Note: The type of dA is ignored for this operation as the entirety of the LAR is sent to the port.