

## # Snow64 Instruction Set

- Notes
  - There are no Instruction LARs (the typical instruction fetch of most computer processors is used instead).
  - Addresses are 64-bit.

- Data LARs:

```
typedef struct packed
{
    // Data field
    union packed
    {
        // This is possibly not valid SystemVerilog because of
        // arrays inside a packed struct, but it makes for nice
        // pseudocode
        logic [7:0] data_8[0:31];

        logic [15:0] data_16[0:15];

        struct packed
        {
            // sign bit, 1 means negative
            logic sign;

            // Exponent, +bias
            logic [7:0] exp;

            // Mantissa; normalized implies 1 MSB
            logic [6:0] mant;
        } data_float_16[0:15];

        logic [31:0] data_32[0:7];
        logic [63:0] data_64[0:3];
    } data;

    // Note that this is a 64-bit structure
    union packed
    {
        struct packed
        {
            logic [63 - 5 : 0] base_ptr;

            logic [4:0] offset;
        } addr_8;

        // Used for both 16-bit integers and the half floats
        struct packed
        {
            logic [63 - 4 : 0] base_ptr;
```

```

        logic [3:0] offset;
    } addr_16;

    struct packed
    {
        logic [63 - 3 : 0] base_ptr;

        logic [2:0] offset;
    } addr_32;

    struct packed
    {
        logic [63 - 2 : 0] base_ptr;

        logic [1:0] offset;
    } addr_64;
} addr;

// Integer Type (used when "is_float_16" is 1'b0):
// 2'b00: 8-bit
// 2'b01: 16-bit
// 2'b10: 32-bit
// 2'b11: 64-bit
logic [1:0] type_of_int;

// These are actually packed 16-bit floats, the implementing the
// high 16 bits of 32-bit IEEE float.
// "type_of_int" is ignored when "is_float_16" is 1'b1.
logic is_float_16;

// Unsigned: 1'b0
// Signed: 1'b1
// Note: unsign_or_sgn is ignored for floats
logic unsign_or_sgn;

// Data should be lazily stored to memory if this is 1'b1
// Otherwise, when this is 1'b0, data in this LAR is up to date
// with memory.
logic dirty;

} DataLar;

```

\* Registers \* The DLARs themselves (there are 16, but this may be changed later): \* dzero (always zero), \* du0, du1, du2, du3 du4, du5, du6, du7, du8, du9, du10, du11 (user registers) \* dlr (standard link register (hardware does not enforce this)) \* dfp (standard frame pointer (hardware does not enforce this)) \* dsp (standard stack pointer (hardware does not enforce this)) \* Other registers: \* pc (the program counter, 64-bit)

# Instruction Set

- Note: All invalid instructions are treated as NOPs.
- ALU Instructions: Opcode Group: 0b000
  - Encoding: 000t aaaa bbbb cccc oooo iiii iiii iiii
    - t: operation type: if 0b0: scalar operation; else: vector operation
    - a: dDest
    - b: dSrc0
    - c: dSrc1
    - o: opcode
    - i: 12-bit signed immediate
  - Note: For ALU instructions, any result that doesn't fit in the destination will be truncated to fit into the destination. This affects both scalar and vector operations.
  - Note: also, for each of these instructions, the address field is not used as an operand, just the data field.
    - Example: `adds d1`
      - Effect: `d1.sdata <= cast_to_type_of_d1(d2.sdata) + cast_to_type_of_d1(d3.sdata)`
  - Instructions:
    - **add** dDest, dSrc0, dSrc1
      - Opcode: 0x0
      - Scalar Mnemonic: `adds`
      - Vector Mnemonic: `addv`
    - **sub** dDest, dSrc0, dSrc1
      - Opcode: 0x1
      - Scalar Mnemonic: `subs`
      - Vector Mnemonic: `subv`
    - **slt** dDest, dSrc0, dSrc1
      - Opcode: 0x2
      - Scalar Mnemonic: `slts`
      - Vector Mnemonic: `sltv`
      - Note: set less than
      - Note: The signedness of dDest will be used for the operation
    - **mul** dDest, dSrc0, dSrc1
      - Opcode: 0x3
      - Scalar Mnemonic: `mul s`
      - Vector Mnemonic: `mul v`
      - Note: If dDest has a larger size than both dSrc0 and dSrc1, then the signedness used for the operation will be that of dDest
      - Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
    - **div** dDest, dSrc0, dSrc1
      - Opcode: 0x4
      - Scalar Mnemonic: `div s`
      - Vector Mnemonic: `div v`

- Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
- **and** dDest, dSrc0, dSrc1
  - Opcode: 0x5
  - Scalar Mnemonic: ands
  - Vector Mnemonic: andv
  - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **orr** dDest, dSrc0, dSrc1
  - Opcode: 0x6
  - Scalar Mnemonic: orrs
  - Vector Mnemonic: orrv
  - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **xor** dDest, dSrc0, dSrc1
  - Opcode: 0x7
  - Scalar Mnemonic: xors
  - Vector Mnemonic: xorv
  - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **shl** dDest, dSrc0, dSrc1
  - Opcode: 0x8
  - Scalar Mnemonic: shls
  - Vector Mnemonic: shlv
  - Note: Shift left
  - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **shr** dDest, dSrc0, dSrc1
  - Opcode: 0x9
  - Scalar Mnemonic: shrs
  - Vector Mnemonic: shrv
  - Note: Shift right
  - Note: dSrc0's signedness is used to determine the type of right shift:
    - If dSrc0 is unsigned, a logical right shift is performed
    - If dSrc0 is signed, an arithmetic right shift is performed
  - Note: dSrc1 is always treated as unsigned (due to being a number of bits to shift by)
  - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **inv** dDest, dSrc0
  - Opcode: 0xa
  - Scalar Mnemonic: invs
  - Vector Mnemonic: invv
  - Note: Bitwise invert
  - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **not** dDest, dSrc0
  - Opcode: 0xb
  - Scalar Mnemonic: nots

- Vector Mnemonic: notv
  - Note: Logical not
- **add** dDest, pc, simm12
  - Opcode: 0xc
  - Scalar Mnemonic: adds
  - Vector Mnemonic: addv
  - Note: This is useful for pc-relative loads, relative branches, and for getting the return address of a subroutine call into a LAR before jumping to a subroutine.
- Instructions for interacting with special-purpose registers:  
Opcode Group: 0b001
  - Encoding: 0010 aaaa oooo iiii iiii iiii iiii iiii
    - a: dA
    - o: opcode
    - i: 20-bit signed immediate
  - Note: all instructions in group 0b001 are scalar operations.
  - Note: dX.sdata is simply the current scalar portion of the data LAR called dX
  - Instructions:
    - **btru** dA, simm20
      - Opcode: 0x0
      - Effect: if (dA.sdata != 0) pc <= pc + sign\_extend\_to\_64(simm20);
    - **bfal** dA, simm20
      - Opcode: 0x1
      - Effect: if (dA.sdata == 0) pc <= pc + sign\_extend\_to\_64(simm20);
    - **jmp** dA
      - Opcode: 0x2
      - Effect: pc <= dA.sdata;
      - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
- Load Instructions: Opcode Group: 0b010
  - Encoding: 0100 aaaa bbbb cccc oooo iiii iiii iiii
    - a: dDest
    - b: dSrc0
    - c: dSrc1
    - o: opcode
    - i: 12-bit signed immediate
  - Effect:
    - Load LAR-sized data from 64-bit address computed as follows: (dB.address + extend\_to\_64(dC.sdata) + (sign\_extend\_to\_64(simm12)))
      - This 64-bit address is referred to as the "effective address".

- The type of extension of the `extend_to_64(dC.sdata)` expression is based upon the type of `dC`.
  - If `dC` is tagged as an unsigned integer, zero-extension is performed.
  - If `dC` is tagged as a signed integer, sign-extension is performed.
  - If `dC` is tagged as a `BFloat16`, `dC.sdata` is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
- Due to associativity of the LARs, these instructions will not actually load from memory if the effective address's data already loaded into a LAR.
- Instructions:
  - **ldu8** `dA, dB, dC, simm12`
    - Opcode: 0x0
    - Note: unsigned 8-bit integer(s)
  - **lds8** `dA, dB, dC, simm12`
    - Opcode: 0x1
    - Note: signed 8-bit integer(s)
  - **ldu16** `dA, dB, dC, simm12`
    - Opcode: 0x2
    - Note: unsigned 16-bit integer(s)
  - **lds16** `dA, dB, dC, simm12`
    - Opcode: 0x3
    - Note: signed 16-bit integer(s)
  - **ldu32** `dA, dB, dC, simm12`
    - Opcode: 0x4
    - Note: unsigned 32-bit integer(s)
  - **lds32** `dA, dB, dC, simm12`
    - Opcode: 0x5
    - Note: signed 32-bit integer(s)
  - **ldu64** `dA, dB, dC, simm12`
    - Opcode: 0x6
    - Note: unsigned 64-bit integer(s)
  - **lds64** `dA, dB, dC, simm12`
    - Opcode: 0x7
    - Note: signed 64-bit integer(s)
  - **ldf16** `dA, dB, dC, simm12`
    - Opcode: 0x8
    - Note: `BFloat16` format floating point number.
- Store Instructions: Opcode Group: 0b011
  - Encoding: 0110 aaaa bbbb cccc oooo iiii iiii iiii
    - a: `dA`
    - b: `dB`
    - c: `dC`
    - o: opcode
    - i: 12-bit signed immediate
  - Note: These are actually type conversion instructions as actual writes to memory are done lazily

- Effect:
  - These instructions mark dA as dirty, change its address to the effective address (see next bullet), and sets its type.
  - The 64-bit effective address is computed as follows:  
 $(dB.address + extend\_to\_64(dC.sdata) + (sign\_extend\_to\_64(simm12)))$
  - The type of extension of the `extend_to_64(dC.sdata)` expression is based upon the type of dC.
    - If dC is tagged as an unsigned integer, zero-extension is performed.
    - If dC is tagged as a signed integer, sign-extension is performed.
    - If dC is tagged as a BFloat16, dC.sdata is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
- Instructions:
  - **stu8** dA, dB, dC, simm12
    - Opcode: 0x0
    - Note: unsigned 8-bit integer(s)
  - **sts8** dA, dB, dC, simm12
    - Opcode: 0x1
    - Note: signed 8-bit integer(s)
  - **stu16** dA, dB, dC, simm12
    - Opcode: 0x2
    - Note: unsigned 16-bit integer(s)
  - **sts16** dA, dB, dC, simm12
    - Opcode: 0x3
    - Note: signed 16-bit integer(s)
  - **stu32** dA, dB, dC, simm12
    - Opcode: 0x4
    - Note: unsigned 32-bit integer(s)
  - **sts32** dA, dB, dC, simm12
    - Opcode: 0x5
    - Note: signed 32-bit integer(s)
  - **stu64** dA, dB, dC, simm12
    - Opcode: 0x6
    - Note: unsigned 64-bit integer(s)
  - **sts64** dA, dB, dC, simm12
    - Opcode: 0x7
    - Note: signed 64-bit integer(s)
  - **stf16** dA, dB, dC, simm12
    - Opcode: 0x8
    - Note: BFloat16 format floating point number.