

Snow64 Instruction Set

- Notes
 - There are no Instruction LARs (the typical instruction fetch of most computer processors is used instead).
 - Addresses are 64-bit.
 - Interrupts are supported (this may be a first for a LARs architecture)
 - Port-mapped I/O is supported (this may be a first for a LARs architecture)
- Data LARs:

```
typedef struct packed
{
    // Data field
    union packed
    {
        // This is possibly not valid SystemVerilog because of
        // arrays inside a packed struct, but it makes for nice
        // pseudocode
        logic [7:0] data_8[0:31];

        logic [15:0] data_16[0:15];

        struct packed
        {
            // sign bit, 1 means negative
            logic sign;

            // Exponent, +bias
            logic [7:0] exp;

            // Mantissa; normalized implies 1 MSB
            logic [6:0] mant;
        } data_float_16[0:15];

        logic [31:0] data_32[0:7];
        logic [63:0] data_64[0:3];
    } data;

    // Note that this is a 64-bit structure
    union packed
    {
        struct packed
        {
            logic [63 - 5 : 0] base_ptr;

            logic [4:0] offset;
        } addr_8;
```

```

    // Used for both 16-bit integers and the half floats
    struct packed
    {
        logic [63 - 4 : 0] base_ptr;

        logic [3:0] offset;
    } addr_16;

    struct packed
    {
        logic [63 - 3 : 0] base_ptr;

        logic [2:0] offset;
    } addr_32;

    struct packed
    {
        logic [63 - 2 : 0] base_ptr;

        logic [1:0] offset;
    } addr_64;
} addr;

// Integer Type (used when "is_float_16" is 1'b0):
// 2'b00: 8-bit
// 2'b01: 16-bit
// 2'b10: 32-bit
// 2'b11: 64-bit
logic [1:0] type_of_int;

// These are actually packed 16-bit floats, the implementing the
// high 16 bits of 32-bit IEEE float.
// "type_of_int" is ignored when "is_float_16" is 1'b1.
logic is_float_16;

// Unsigned: 1'b0
// Signed: 1'b1
// Note: unsgn_or_sgn is ignored for floats
logic unsgn_or_sgn;

// Data should be lazily stored to memory if this is 1'b1
// Otherwise, when this is 1'b0, data in this LAR is up to date
// with memory.
logic dirty;

} DataLar;

```

* Registers * The DLARs themselves (there are 16, but this may be changed later): * dzero (always zero), * du0, du1, du2, du3, du4, du5, du6, du7, du8, du9, du10, du11 (user registers) * dlr (standard link register (hardware does not enforce this)) * dfp (standard frame pointer (hardware does not enforce

this)) * dsp (standard stack pointer (hardware does not enforce this)) *
 Other registers: * pc (the program counter, 64-bit) * ie (whether or not
 interrupts are enabled, 1-bit) * ireta (the interrupt return address, 64-bit) *
 idsta (the interrupt destination address, 64-bit; upon an interrupt, the
 program counter is set to the value in this register)

Instruction Set

- Note: All invalid instructions are treated as NOPs.
- ALU Instructions: Opcode Group: 0b000
 - Encoding: 000t aaaa bbbb cccc oooo iiii iiii iiii
 - t: operation type: if 0b0: scalar operation; else: vector operation
 - a: dDest
 - b: dSrc0
 - c: dSrc1
 - o: opcode
 - i: 12-bit signed immediate
 - Note: For ALU instructions, any result that doesn't fit in the destination will be truncated to fit into the destination. This affects both scalar and vector operations.
 - Note: also, for each of these instructions, the address field is not used as an operand, just the data field.
 - Example: adds d1, d2, d3
 - Effect: $d1.sdata \leq \text{cast_to_type_of_d1}(d2.sdata) + \text{cast_to_type_of_d1}(d3.sdata)$
 - Instructions:
 - **add** dDest, dSrc0, dSrc1
 - Opcode: 0x0
 - Scalar Mnemonic: adds
 - Vector Mnemonic: addv
 - **sub** dDest, dSrc0, dSrc1
 - Opcode: 0x1
 - Scalar Mnemonic: subs
 - Vector Mnemonic: subv
 - **slt** dDest, dSrc0, dSrc1
 - Opcode: 0x2
 - Scalar Mnemonic: slts
 - Vector Mnemonic: sltv
 - Note: set less than
 - Note: The signedness of dDest will be used for the operation
 - **mul** dDest, dSrc0, dSrc1
 - Opcode: 0x3
 - Scalar Mnemonic: muls
 - Vector Mnemonic: mulv
 - Note: If dDest has a larger size than both dSrc0 and dSrc1, then the signedness used for the operation will be that of dDest

- Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
- **div** dDest, dSrc0, dSrc1
 - Opcode: 0x4
 - Scalar Mnemonic: divs
 - Vector Mnemonic: divv
 - Note: This operation is not guaranteed to be single cycle, and thus pipeline stalls will be used
- **and** dDest, dSrc0, dSrc1
 - Opcode: 0x5
 - Scalar Mnemonic: ands
 - Vector Mnemonic: andv
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **orr** dDest, dSrc0, dSrc1
 - Opcode: 0x6
 - Scalar Mnemonic: orrs
 - Vector Mnemonic: orrv
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **xor** dDest, dSrc0, dSrc1
 - Opcode: 0x7
 - Scalar Mnemonic: xors
 - Vector Mnemonic: xorv
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **shl** dDest, dSrc0, dSrc1
 - Opcode: 0x8
 - Scalar Mnemonic: shls
 - Vector Mnemonic: shlv
 - Note: Shift left
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **shr** dDest, dSrc0, dSrc1
 - Opcode: 0x9
 - Scalar Mnemonic: shrs
 - Vector Mnemonic: shrv
 - Note: Shift right
 - Note: dSrc0's signedness is used to determine the type of right shift:
 - If dSrc0 is unsigned, a logical right shift is performed
 - If dSrc0 is signed, an arithmetic right shift is performed
 - Note: dSrc1 is always treated as unsigned (due to being a number of bits to shift by)
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
- **inv** dDest, dSrc0
 - Opcode: 0xa
 - Scalar Mnemonic: invs
 - Vector Mnemonic: invv

- Note: Bitwise invert
 - Note: For floats, this operation treats all operands as 16-bit signed integers.
 - **not** dDest, dSrc0
 - Opcode: 0xb
 - Scalar Mnemonic: nots
 - Vector Mnemonic: notv
 - Note: Logical not
 - **add** dDest, pc, simm12
 - Opcode: 0xc
 - Scalar Mnemonic: adds
 - Vector Mnemonic: addv
 - Note: This is useful for pc-relative loads, relative branches, and for getting the return address of a subroutine call into a LAR before jumping to a subroutine.
- Instructions for interacting with special-purpose registers:
 Opcode Group: 0b001
 - Encoding: 0010 aaaa oooo iiii iiii iiii iiii iiii
 - a: dA
 - o: opcode
 - i: 20-bit signed immediate
 - Note: all instructions in group 0b001 are scalar operations.
 - Note: dX.sdata is simply the current scalar portion of the data LAR called dX
 - Instructions:
 - **btru** dA, simm20
 - Opcode: 0x0
 - Effect: if (dA.sdata != 0) pc <= pc + sign_extend_to_64(simm20);
 - **bfal** dA, simm20
 - Opcode: 0x1
 - Effect: if (dA.sdata == 0) pc <= pc + sign_extend_to_64(simm20);
 - **jmp** dA
 - Opcode: 0x2
 - Effect: pc <= dA.sdata;
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **ei**
 - Opcode: 0x3
 - Effect: ie <= 1'b1;
 - Note: Enable interrupts
 - **di**
 - Opcode: 0x4
 - Effect: ie <= 1'b0;
 - Note: Disable interrupts
 - **reti**
 - Opcode: 0x5
 - Effect: ie <= 1'b1; pc <= ireta;

- Note: Return from an interrupt
 - **cpy dA, ie**
 - Opcode: 0x6
 - Effect: `dA.sdata <= ie;` // acts differently if dA is tagged as a float
 - **cpy dA, ireta**
 - Opcode: 0x7
 - Effect: `dA.sdata <= ireta;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **cpy dA, idsta**
 - Opcode: 0x8
 - Effect: `dA.sdata <= idsta;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **cpy ie, dA**
 - Opcode: 0x9
 - Effect: `ie <= (dA.sdata != 0);`
 - **cpy ireta, dA**
 - Opcode: 0xa
 - Effect: `ireta <= dA.sdata;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
 - **cpy idsta, dA**
 - Opcode: 0xb
 - Effect: `idsta <= dA.sdata;`
 - Note: It is suggested to have dA.sdata be at least as large as the largest memory address (which might not be 64-bit if there isn't enough physical memory for that)
- Load Instructions: Opcode Group: 0b010
 - Encoding: 0100 aaaa bbbb cccc oooo iiii iiii iiii
 - a: dDest
 - b: dSrc0
 - c: dSrc1
 - o: opcode
 - i: 12-bit signed immediate
 - Effect:
 - Load LAR-sized data from 64-bit address computed as follows: `(dB.address + extend_to_64(dC.sdata) + (sign_extend_to_64(simm12)))`
 - This 64-bit address is referred to as the "effective address".
 - The type of extension of the `extend_to_64(dC.sdata)` expression is based upon the type of dC.
 - If dC is tagged as an unsigned integer, zero-extension is performed.
 - If dC is tagged as a signed integer, sign-extension is performed.

- If dC is tagged as a BFloat16, dC.sdata is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
 - Due to associativity of the LARs, these instructions will not actually load from memory if the effective address's data already loaded into a LAR.
- Instructions:
 - **ldu8** dA, dB, dC, simm12
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **lds8** dA, dB, dC, simm12
 - Opcode: 0x1
 - Note: signed 8-bit integer(s)
 - **ldu16** dA, dB, dC, simm12
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
 - **lds16** dA, dB, dC, simm12
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
 - **ldu32** dA, dB, dC, simm12
 - Opcode: 0x4
 - Note: unsigned 32-bit integer(s)
 - **lds32** dA, dB, dC, simm12
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
 - **ldu64** dA, dB, dC, simm12
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
 - **lds64** dA, dB, dC, simm12
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
 - **ldf16** dA, dB, dC, simm12
 - Opcode: 0x8
 - Note: BFloat16 format floating point number.
- Store Instructions: Opcode Group: 0b011
 - Encoding: 0110 aaaa bbbb cccc oooo iiii iiii iiii
 - a: dA
 - b: dB
 - c: dC
 - o: opcode
 - i: 12-bit signed immediate
 - Note: These are actually type conversion instructions as actual writes to memory are done lazily
 - Effect:
 - These instructions mark dA as dirty, change its address to the effective address (see next bullet), and sets its type.
 - The 64-bit effective address is computed as follows:
 $(dB.address + \text{extend_to_64}(dC.sdata) + (\text{sign_extend_to_64}(\text{simm12})))$

- The type of extension of the `extend_to_64(dC.sdata)` expression is based upon the type of `dC`.
 - If `dC` is tagged as an unsigned integer, zero-extension is performed.
 - If `dC` is tagged as a signed integer, sign-extension is performed.
 - If `dC` is tagged as a `BFloat16`, `dC.sdata` is casted to a 64-bit signed integer. (This one is weird... normally, addressing isn't done with floating point numbers!).
- Instructions:
 - **stu8** `dA, dB, dC, simm12`
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **sts8** `dA, dB, dC, simm12`
 - Opcode: 0x1
 - Note: signed 8-bit integer(s)
 - **stu16** `dA, dB, dC, simm12`
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
 - **sts16** `dA, dB, dC, simm12`
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
 - **stu32** `dA, dB, dC, simm12`
 - Opcode: 0x4
 - Note: unsigned 32-bit integer(s)
 - **sts32** `dA, dB, dC, simm12`
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
 - **stu64** `dA, dB, dC, simm12`
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
 - **sts64** `dA, dB, dC, simm12`
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
 - **stf16** `dA, dB, dC, simm12`
 - Opcode: 0x8
 - Note: `BFloat16` format floating point number.
- Port-mapped Input/Output Instructions: Opcode Group: 0b100
 - Encoding: 100t aaaa bbbb oooo iiii iiii iiii iiii
 - t: operation type: if 0b0: scalar operation; else: vector operation
 - a: `dA`
 - b: `dB`
 - o: opcode
 - i: 16-bit signed immediate
 - Note: `dX.sdata` is simply the current scalar portion of the data LAR called `dX`
 - Note: For the `in...` instructions, the entirety of `dA.data` is set to the received data. The type of `dA` is set based upon the instruction opcode.

- Note: For `outs`, `dA.sdata` is sent to the output port, along with the type of data (in case the particular I/O port cares).
- Note: For `outv`, the entirety of `dA.data` is sent to the output port, along with the type of data (in case the particular I/O port cares).
- Note: For each of these instructions, the I/O address used is computed by the formula `cast_to_64(dB.sdata) + sign_extend_to_64(simm16)`
- Instructions:
 - **inu8** `dA, dB, simm16`
 - Opcode: 0x0
 - Note: unsigned 8-bit integer(s)
 - **ins8** `dA, dB, simm16`
 - Opcode: 0x1
 - Note: signed 8-bit integer(s)
 - **inu16** `dA, dB, simm16`
 - Opcode: 0x2
 - Note: unsigned 16-bit integer(s)
 - **ins16** `dA, dB, simm16`
 - Opcode: 0x3
 - Note: signed 16-bit integer(s)
 - **inu32** `dA, dB, simm16`
 - Opcode: 0x4
 - Note: unsigned 32-bit integer(s)
 - **ins32** `dA, dB, simm16`
 - Opcode: 0x5
 - Note: signed 32-bit integer(s)
 - **inu64** `dA, dB, simm16`
 - Opcode: 0x6
 - Note: unsigned 64-bit integer(s)
 - **ins64** `dA, dB, simm16`
 - Opcode: 0x7
 - Note: signed 64-bit integer(s)
 - **inf16** `dA, dB, simm16`
 - Opcode: 0x8
 - Note: BFloat16 format floating point number.
 - **out** (actual mnemonics below)
 - Opcode: 0x9
 - **outs** `dA, dB, simm16`
 - t: 0
 - Note: `dA.sdata` is simply sent to the output data bus.
 - **outv** `dA, dB, simm16`
 - t: 1
 - Note: The type of `dA` is ignored for this operation as the entirety of the LAR is sent to the port.